

Assignment 4: Model-Based RL and Exploration

Andrew ID: mukaiy

NOTE: Please do **NOT** change the sizes of the answer blocks or plots.

1 Problem 1: Dynamics Model Training – [10 points total]

Theory questions

The third model performs the best, because it achieves the least $MPE = 0.07804489$.
More training steps per iteration improves convergence a lot, and larger MLP interpolates better.

Plot




Figure 1: 500 training steps per iteration, 1 x 32 MLP.

Plot




Figure 2: 5 training steps per iteration, 2 x 250 MLP.

Plot




Figure 3: 500 training steps per iteration, 2 x 250 MLP.

2 Problem 2: Action Selection

Plot

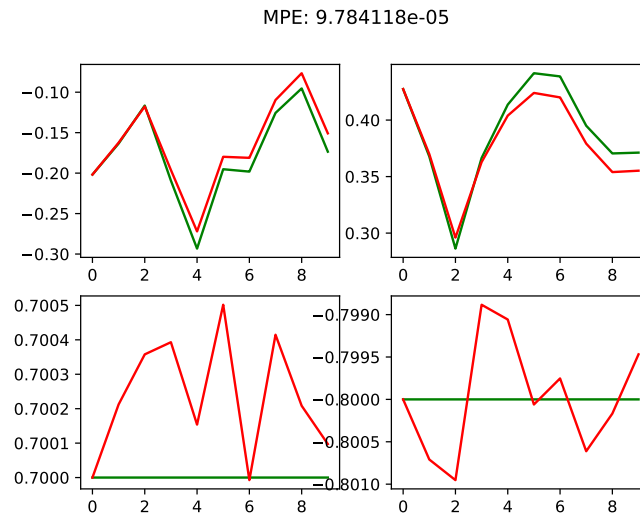


Figure 4: Optimizing the Dynamics Model where MPC needs to be used improves a lot.

3 Problem 3: Iterative Model Training

Plot




Figure 5: Average evaluation return over training iterations.

Orange: cheetah. Purple: obstacles. Green: reacher.

Plot




Figure 6: Standard deviation evaluation return over training iterations.

Orange: cheetah. Purple: obstacles. Green: reacher.

Plot





Figure 7: Average training return over training iterations.


Orange: cheetah. Purple: obstacles. Green: reacher.


4 Problem 4: Hyper-parameter Comparison

Plot




Plot






5 Problem 5: Hyper-parameter Comparison (Bonus)




6 Problem 6: Exploration (Bonus)


Plot



Plot



Plot



Plot

